

GEMM Analytical Model - Equations

Auto-generated from PYTHON_FORMULAS and LATEX_SYMBOL_OVERRIDES in gemmanalytics.py.

Output order: execution

row 22 - INPUT_A_BYTES_PER_ELEMENT

$$\text{Bytes}_{\text{pElement}}^{(\text{matA})} = \text{DataFormatToBytes}[\text{Fmt}^{(\text{matA})}]$$

row 23 - INPUT_B_BYTES_PER_ELEMENT

$$\text{Bytes}_{\text{pElement}}^{(\text{matB})} = \text{DataFormatToBytes}[\text{Fmt}^{(\text{matB})}]$$

row 25 - OUTPUT_BYTES_PER_ELEMENT_AFTER_DOWN_CONVERSION

$$\text{Bytes}_{\text{pElement}}^{(\text{D}\downarrow)} = \text{DataFormatToBytes}[\text{Fmt}^{(\text{matD}\downarrow)}]$$

row 33 - EU_COUNT

$$|\text{EU}| = |\text{XeCore}|_{\text{pXeCU}} \times |\text{EU}|_{\text{pXeCore}} \times |\text{XeCU}|$$

row 102 - TOTAL_L2_SIZE_B_FOR_A_SINGLE_INSTANCE

$$|\text{L2Bytes}| = \text{L2BankSize}_{\text{MB}} \times |\text{L2Banks}|_{\text{pXeCU}} \times 1024 \times 1024$$

row 125 - MAX_POSSIBLE_HBM_BW_FREQ_B_CLK

$$\text{MemBW}_{\text{BpClk}}^{\text{max}} = \frac{\text{MemBW}_{\text{GBps}}^{\text{max}}}{f_{\text{GHz}}^{(\text{GT})}}$$

row 36 - MMA_MAC_THROUGHPUT_PER_XECORE

$$\tau_{\text{mMACpClk-XeCore}}^{(\text{peak})} = \min\left(\frac{4}{\frac{\text{FLOOR}(\text{Bytes}_{\text{pElement}}^{(\text{matA})} \times 2, 1)}{2}}, \frac{4}{\frac{\text{FLOOR}(\text{Bytes}_{\text{pElement}}^{(\text{matB})} \times 2, 1)}{2}}\right) \times \text{D}_{\text{DPAS}} \times 16 \times |\text{EU}|_{\text{pXeCore}}$$

row 42 - CLKS_PER_DPAS

$$\text{CLKS}_{\text{DPAS}} = \frac{\text{M}_{\text{pThread}} \times \text{K}_{\text{pThread}} \times \text{N}_{\text{pThread}}}{\frac{\tau_{\text{mMACpClk-XeCore}}^{(\text{peak})}}{|\text{EU}|_{\text{pXeCore}}}}$$

row 13 - WAVES

$$\text{N}_{\text{waves}} = \text{CEILING}(|\text{Tiles}|_{\text{N}}^{(\text{GPU})}, 1) \times \text{CEILING}(|\text{Tiles}|_{\text{M}}^{(\text{GPU})}, 1)$$

row 14 - TG_TILES_IN_N

$$|\text{Tiles}|_{\text{N}}^{(\text{TG})} = \frac{\text{N}_{\text{dim}}}{\text{W}_{\text{element}}^{(\text{ThreadGroup, realized by multiple MMA iterations})}}$$

row 15 - TG_TILES_IN_M

$$|\text{Tiles}|_{\text{M}}^{(\text{TG})} = \frac{\text{M}_{\text{dim}}}{\text{H}_{\text{element}}^{(\text{ThreadGroup})}}$$

row 16 - TG_CLUSTER_TILES_IN_N

$$|\text{Tiles}|_N^{(\text{TG Cluster})} = \frac{N_d \text{im}}{W_{\text{element}}^{(\text{XeCoreCluster})}}$$

row 17 - TG_CLUSTER_TILES_IN_M

$$|\text{Tiles}|_M^{(\text{TG Cluster})} = \frac{M_d \text{im}}{H_{\text{element}}^{(\text{XeCoreCluster})}}$$

row 18 - XECU_TILES_IN_N

$$|\text{Tiles}|_N^{(\text{XeCU})} = \frac{N_d \text{im}}{W_{\text{element}}^{(\text{XeCUTile})}}$$

row 19 - XECU_TILES_IN_M

$$|\text{Tiles}|_M^{(\text{XeCU})} = \frac{M_d \text{im}}{H_{\text{element}}^{(\text{XeCUTile})}}$$

row 20 - GPU_TILES_IN_N

$$|\text{Tiles}|_N^{(\text{GPU})} = \frac{\frac{N_d \text{im}}{W_{\text{element}}^{(\text{XeCUTile})}}}{W_{\text{XeCU}}^{(\text{GPUTile})}}$$

row 21 - GPU_TILES_IN_M

$$|\text{Tiles}|_M^{(\text{GPU})} = \frac{\frac{M_d \text{im}}{H_{\text{element}}^{(\text{XeCUTile})}}}{H_{\text{XeCU}}^{(\text{GPUTile})}}$$

row 43 - THREAD_WIDTH_IN_UNITS_OF_ELEMENTS

$$W_{\text{element}}^{(\text{Thread})} = N_{\text{pThread}}$$

row 44 - THREAD_HEIGHT_IN_UNITS_OF_ELEMENTS

$$H_{\text{element}}^{(\text{Thread})} = M_{\text{pThread}}$$

row 47 - TG_WIDTH_IN_UNITS_OF_ELEMENT_REALIZED_BY_MULTIPLE_MMA_ITERATIONS

$$W_{\text{element}}^{(\text{ThreadGroup, realized by multiple MMA iterations})} = W_{\text{thread}}^{(\text{ThreadGroup})} \times W_{\text{element}}^{(\text{Thread})}$$

row 48 - TG_HEIGHT_IN_UNITS_OF_ELEMENT

$$H_{\text{element}}^{(\text{ThreadGroup})} = H_{\text{thread}}^{(\text{ThreadGroup})} \times H_{\text{element}}^{(\text{Thread})}$$

row 52 - XECORE_CLUSTER_WIDTH_IN_UNITS_OF_ELEMENT

$$W_{\text{element}}^{(\text{XeCoreCluster})} = W_{\text{TG, keep cluster sizes as 4}}^{(\text{XeCoreCluster})} \times W_{\text{element}}^{(\text{ThreadGroup, realized by multiple MMA iterations})}$$

row 53 - XECORE_CLUSTER_HEIGHT_IN_UNITS_OF_ELEMENT

$$H_{\text{element}}^{(\text{XeCoreCluster})} = H_{\text{TG, keep cluster sizes as 4}}^{(\text{XeCoreCluster})} \times H_{\text{element}}^{(\text{ThreadGroup})}$$

row 56 - XECU_TILE_WIDTH_IN_UNITS_OF_ELEMENT

$$W_{\text{element}}^{(\text{XeCUTile})} = W_{\text{TG}}^{(\text{XECUTile})} \times W_{\text{element}}^{(\text{ThreadGroup, realized by multiple MMA iterations})}$$

row 57 - XECU_TILE_HEIGHT_IN_UNITS_OF_ELEMENT

$$H_{\text{element}}^{(\text{XeCUTile})} = H_{\text{TG}}^{(\text{XECUTile})} \times H_{\text{element}}^{(\text{ThreadGroup})}$$

row 59 - GPU_TILE_HEIGHT_IN_XECU_UINT

$$H_{\text{XeCU}}^{(\text{GPUTile})} = \frac{|\text{XeCU}|}{W_{\text{XeCU}}^{(\text{GPUTile})}}$$

row 60 - GPU_TILE_WIDTH_IN_UNITS_OF_ELEMETNS

$$W_{\text{element}}^{(\text{GPUTile})} = W_{\text{XeCU}}^{(\text{GPUTile})} \times W_{\text{element}}^{(\text{XeCUTile})}$$

row 61 - GPU_TILE_HEIGHT_IN_UNITS_OF_ELEMETNS

$$H_{\text{element}}^{(\text{GPUTile})} = H_{\text{XeCU}}^{(\text{GPUTile})} \times H_{\text{element}}^{(\text{XeCUTile})}$$

row 63 - MAT_A_INPUT_SIZE_B

$$\text{Size}_{\text{B}}^{(\text{matA})} = M_{\text{dim}} \times K_{\text{dim}} \times \text{Bytes}_{\text{pElement}}^{(\text{matA})}$$

row 64 - MAT_B_INPUT_SIZE_B

$$\text{Size}_{\text{B}}^{(\text{matB})} = K_{\text{dim}} \times N_{\text{dim}} \times \text{Bytes}_{\text{pElement}}^{(\text{matB})}$$

row 65 - MAT_C_INPUT_D_OUTPUT_SIZE_B

$$\text{Size}_{\text{matB}}^{(\text{matC, matD})} = M_{\text{dim}} \times N_{\text{dim}} \times \text{Bytes}_{\text{pElement}}^{(\text{D}\downarrow)}$$

row 66 - MAT_D_INTERMEDIATE_SIZE_B

$$\text{Size}_{\text{B}}^{(\text{matD}\downarrow)} = M_{\text{dim}} \times N_{\text{dim}} \times \text{Byte}_{\text{pElement}}^{(\text{matD})}$$

row 103 - WORKING_DATA_SET_SIZE_OF_K_IN_L2_CORRESP_20K_CLOCKS_OF_THREAD_DIVERGENCE

$$|\text{WorkingSet}|_{20\text{Kclksofthreaddivergence}}^{(\text{KinL2})} = \min(20000 \times \frac{K_{\text{pThread}}}{\text{CLKS}_{\text{DPAS}}}, K_{\text{dim}})$$

row 104 - TOTAL_REQUIRED_L2_SIZE_FOR_IDEAL_HIT_RATE_B_FOR_A_SINGLE_INSTANCE_AND_SINGLE_WAVE

$$\text{L2Rd}_{\text{B}}^{(\text{total})} = H_{\text{element}}^{(\text{XeCUTile})} \times \text{Bytes}_{\text{pElement}}^{(\text{matA})} \times |\text{WorkingSet}|_{20\text{Kclksofthreaddivergence}}^{(\text{KinL2})} + W_{\text{element}}^{(\text{XeCUTile})} \times \text{Bytes}_{\text{pElement}}^{(\text{matB})} \times |\text{WorkingSet}|_{20\text{Kclksofthreaddivergence}}^{(\text{KinL2})} + W_{\text{element}}^{(\text{XeCUTile})} \times H_{\text{element}}^{(\text{XeCUTile})}$$

row 37 - CLK_SPECIFIED EFFICIENCY

$$T_{\text{clk}}^{(\text{total})} = \frac{\frac{\frac{M_{\text{dim}} \times K_{\text{dim}} \times N_{\text{dim}}}{|\text{XeCore}|_{\text{pXeCU}}}}{|\text{XeCU}|}}{\tau_{\text{mMACpClk-XeCore}}^{(\text{peak})} \eta_{\text{systolic}}}$$

row 69 - TOTAL_L2_READ_B

$$\text{L2Rd}_{\text{B}}^{(\text{total})} = \text{Size}_{\text{B}}^{(\text{matA})} \times \text{CEILING}(\frac{N_{\text{dim}}}{W_{\text{element}}^{(\text{ThreadGroup, realized by multiple MMA iterations})}}, 1) + \text{Size}_{\text{B}}^{(\text{matB})} \times \text{CEILING}(\frac{M_{\text{dim}}}{H_{\text{element}}^{(\text{ThreadGroup})}}, 1)$$

row 70 - TOTAL_L2_WRITE_B

$$L2Wr_B^{(total)} = Size_{matB}^{(matC, matD)}$$

row 71 - TOTAL_L1_READ_B

$$L1Rd_B^{(total)} = Size_B^{(matA)} \times CEILING(\frac{N_{dim}}{W_{element}^{(Thread)}}, 1) + Size_B^{(matB)} \times CEILING(\frac{M_{dim}}{H_{element}^{(Thread)}}, 1)$$

row 72 - TOTAL_L1_WRITE_B

$$L1Wr_B^{(total)} = Size_{matB}^{(matC, matD)}$$

row 74 - L2_READ_B_XECORE_CLK

$$L2RdB_{p(Clk \cdot XeCore)} = \frac{\frac{L2Rd_B^{(total)}}{|XeCU| \times |XeCore|_{pXeCU}}}{T_{clk}^{(total)}}$$

row 75 - L2_WRITE_B_XECORE_CLK

$$L2WrB_{p(Clk \cdot XeCore)} = \frac{\frac{L2Wr_B^{(total)}}{|XeCU| \times |XeCore|_{pXeCU}}}{T_{clk}^{(total)}}$$

row 76 - L2_READ_WRITE_B_XECORE_CLK

$$L2RdWrB_{p(Clk \cdot XeCore)} = L2RdB_{p(Clk \cdot XeCore)} + L2WrB_{p(Clk \cdot XeCore)}$$

row 77 - L1_READ_B_XECORE_CLK

$$L1RdB_{p(Clk \cdot XeCore)} = L1RdB_{p(Clk \cdot EU)} \times |EU|_{pXeCore}$$

row 78 - L1_WRITE_B_XECORECLK

$$L1WrB_{p(Clk \cdot XeCore)} = L1WrB_{p(Clk \cdot EU)} \times |EU|_{pXeCore}$$

row 79 - L1_READ_B_EU_CLK

$$L1RdB_{p(Clk \cdot EU)} = \frac{\frac{L1Rd_B^{(total)}}{|EU|}}{T_{clk}^{(total)}}$$

row 80 - L1_WRITE_B_EU_CLK

$$L1WrB_{p(Clk \cdot EU)} = \frac{\frac{L1Wr_B^{(total)}}{|EU|}}{T_{clk}^{(total)}}$$

row 83 - L2_READ_MAX_B_XECORE_CLK

$$L2RdB_{p(Clk \cdot XeCore)}^{(max)} = \frac{|XeCore|_{pXeCU} \times 64}{|XeCore|_{pXeCU}}$$

row 84 - L2_WRITE_MAX_B_XECORE_CLK

$$L2WrB_{p(Clk \cdot XeCore)}^{(max)} = L2RdB_{p(Clk \cdot XeCore)}^{(max)}$$

row 85 - L2_READ_WRITE_MAX_B_XECORE_CLK

$$L2RdWrB_{p(Clk.XeCore)}^{(max)} = L2WrB_{p(Clk.XeCore)}^{(max)}$$

row 90 - L2_READ_B_XECORE_CLK_PCT

$$\eta_{L2RdB/p(Clk.XeCore)} = \frac{L2RdB_{p(Clk.XeCore)}}{L2RdB_{p(Clk.XeCore)}^{(max)}}$$

row 91 - L2_WRITE_B_XECORE_CLK_PCT

$$\eta_{L2WrB/p(Clk.XeCore)} = \frac{L2WrB_{p(Clk.XeCore)}}{L2WrB_{p(Clk.XeCore)}^{(max)}}$$

row 92 - L2_READ_WRITE_B_XECORE_CLK_PCT

$$\eta_{L2RdWrB/p(Clk.XeCore)} = \frac{L2RdWrB_{p(Clk.XeCore)}}{L2RdWrB_{p(Clk.XeCore)}^{(max)}}$$

row 93 - L1_READ_B_EU_CLK_PCT

$$\eta_{L1RdB/p(Clk.EU)} = \frac{L1RdB_{p(Clk.EU)}}{L1RdBW_{BpClk.EU}^{max}}$$

row 94 - L1_WRITE_B_EU_CLK_PCT

$$\eta_{L1WrB/p(Clk.EU)} = \frac{L1WrB_{p(Clk.EU)}}{L1WrBW_{BpClk.EU}^{max}}$$

row 106 - L2_HIT_RATE_ASSUMED_RANDOM_ACCESS_WITHIN_THE_WORKING_DATA_SET_PCT

$$L2HitRate_{random\ WS\ access} = \min\left(\frac{|L2Bytes|}{TotalRequiredL2Size_{matB\ alwayshit}^{1\ instance,\ 1\ wave}}, 1\right)$$

row 107 - L2_MISS_RATE_PCT

$$L2MissRate = 1 - L2HitRate_{random\ WS\ access}$$

row 108 - TOTAL_L2_READ_TRAFFIC_B

$$L2RdB_{total} = L2RdB_B^{(total)}$$

row 110 - PROBABILITY_OF_MATA_HIT_IN_L2_DURING_A_NON_FIRST_WAVE_PCT

$$P_{L2Hit,after1stwave}^{(matB)} = IF(Size_B^{(matA)} + Size_B^{(matB)} + Size_{matB}^{(matC,matD)} \leq |L2Bytes|, 1.0, IF(K_{dim} > 2 \times |WorkingSet|_{20Kclksofthreaddivergence}^{(KinL2)}, 0.0, 1 - \frac{K_{dim} - |WorkingSet|_{20Kclksofthreaddivergence}^{(KinL2)}}{|WorkingSet|_{20Kclksofthreaddivergence}^{(KinL2)}})$$

row 111 - PROBABILITY_OF_MATB_HIT_IN_L2_DURING_A_NON_FIRST_WAVE_PCT

$$P_{L2Hit,after1stwave}^{(matB)} = IF(Size_B^{(matA)} + Size_B^{(matB)} + Size_{matB}^{(matC,matD)} \leq |L2Bytes|, 1.0, 0.0)$$

row

112

- PROBABILITY_OF_MATA_MISS_IN_L2_DURING_A_NON_FIRST_WAVE_PCT

$$P_{L2Miss,after1stwave}^{(matA)} = 1 - P_{L2Hit,after1stwave}^{(matA)}$$

row

113

- PROBABILITY_OF_MATB_MISS_IN_L2_DURING_A_NON_FIRST_WAVE_PCT

$$P_{L2Miss,after1stwave}^{(matB)} = 1 - P_{L2Hit,after1stwave}^{(matB)}$$

row 116 - TOTAL_HBM_READ_B_AFTER_A_COMPLETION_OF_A_WAVE_CONSIDER_COLD_CACHE

$$Size_B^{(matB)} + Size_B^{(matB)} \times \left(\text{CEILING}\left(\frac{M_{\text{element}}^{im}}{H_{\text{element}}^{XCUTile}}, 1\right) - 1 \right) \times P_{L2Miss,after1stwave}^{(matB)} + \left(L2RdB_{total} - \left(Size_B^{(matA)} + Size_B^{(matA)} \times \left(\text{CEILING}\left(\frac{N_{\text{element}}^{im}}{W_{\text{element}}^{XCUTile}}, 1\right) - 1 \right) \times P_{L2Miss,after1stwave}^{(matA)} \right) \right) \times P_{L2Miss,after1stwave}^{(matA)}$$

2

row 117 - TOTAL_HBM_WRITE_B

$$MemWrB^{(total)} = Size_{matB}^{(matC,matD)}$$

row 118 - TOTAL_HBM_B

$$MemRdWrB^{(total)} = MemWrB^{(total)} + MemRdB_{coldstart}^{(wave)}$$

row 121 - HBM_READ_B_CLK

$$MemRdB_{pClk} = \frac{MemRdB_{coldstart}^{(wave)}}{T_{clk}^{(total)}}$$

row 122 - HBM_WRITE_B_CLK

$$MemWrB_{pClk} = \frac{MemWrB^{(total)}}{T_{clk}^{(total)}}$$

row 123 - HBM_TOTAL_B_CLK

$$MemRdWrB_{pClk} = \frac{MemRdWrB^{(total)}}{T_{clk}^{(total)}}$$

row 126 - HBM_BW_PCT

$$\eta_{MemBW} = \frac{MemRdWrB_{pClk}}{MemBW_{BpClk}^{max}}$$

row 127 - GTI_READ_BW_PCT

$$\eta_{GTIRdBW} = \frac{MemRdB_{pClk}}{GtiRdBW_{BpClk}^{max}}$$

row 128 - GTI_WRITE_BW_PCT

$$\eta_{GTIWrBW} = \frac{MemWrB_{pClk}}{GtiWrBW_{BpClk}^{max}}$$